



Carnegie Mellon University
Language Technologies Institute



Complex Few-shot Reasoning with LLMs

Aman Madaan

Complex Few-shot Reasoning with LLMs

Goal Today:

- Try to cover the details of all possible exciting recent works 
- Provide a high-level summary of a *class* of few-shot prompting techniques 

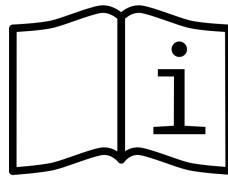
I will definitely miss interesting works:

- Please reach out! We will add it to our website

Complex Few-shot Reasoning with LLMs: Key Techniques



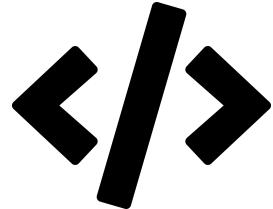
Reasoning Elaboration




Instructions



Tool Augmentation



Structured Generation



Feedback



Memory

Background

Solving middle-school math problems

Q: *Shawn has 5 toys. For Christmas, he got 2 toys each from his mom and dad. How many toys does he have now?*



+



=




A: The answer is 9 toys




Fine-tuning

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?



A: The answer is 5 cars.


Q: Leah had 32 chocolates and her sister had 42. If they ate 35, how many pieces do they have left in total?



A: The answer is 39 pieces.

Train/Fine-tune

Q: Shawn has five toys. For Christmas, he got two toys each from his mom and dad. How many toys does he have now?



A: The answer is 9 toys


Test

Few-shot prompting

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?

A: *The answer is 5 cars.*

Prompt



The answer is 9 toys

Q: Shawn has five toys. For Christmas, he got two toys each from his mom and dad. How many toys does he have now?

A:

Test Example


Rest of the slides

- General idea of the technique
- Representative Work

Complex Few-shot Reasoning with LLMs: Key Techniques




Reasoning Elaboration



Instructions



Tool Augmentation



Structured Generation




Feedback




Memory

Complex Few-shot Reasoning with LLMs: Key Techniques




Techniques overlap in practice – focus on the main contribution

Complex Few-shot Reasoning with LLMs: Key Techniques




Reasoning Elaboration




Instructions



Tool Augmentation



Structured Generation



Feedback




Memory

Few-shot prompting

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?

A: *The answer is 5 cars.*

Prompt



The answer is 9 toys

Q: Shawn has five toys. For Christmas, he got two toys each from his mom and dad. How many toys does he have now?

A:

Test Example

Few-shot prompting

*Direct
Prompt*

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?

A: *The answer is 5 cars.*

Chain-of-thought Prompting (Wei et al. 2022)

*Direct
Prompt*

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?

A: *The answer is 5 cars.*

*Chain-of-Thought
Prompt*


Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?

Thought (T): There are originally 3 cars. 2 more cars arrive. $3 + 2 = 5$.

A: *The answer is 5 cars.*

Chain of thought prompting is extremely effective

PaLM 62B



Chain-of-thought Prompting (CoT)

- General idea:
 - Standard prompt:
 - $Q \rightarrow A$
 - Chain-of-thought prompt:
 - $Q \rightarrow \text{Reasoning Process, } A$
- Many existing prompting techniques can be seen as an improvement over the general CoT strategy
 - The thought can be: Text, Code, API calls
- Similar ideas:
 - [Program Induction by Rationale Generation: Learning to Solve and Explain Algebraic Word Problems \(Ling et al. 2017\)](#)
 - [Think about it! Improving defeasible reasoning by first modeling the question scenario \(Madaan et al. 2021\)](#)
 - [Show your work: Scratchpads for Intermediate Computation with Language Models \(Nye et al. 2021\)](#)

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?

Thought (T): There are originally 3 cars. 2 more cars arrive. $3 + 2 = 5$.

A: The answer is 5 cars.

Least-to-Most Prompting (Zhou et al. 2022)

- Breakdown the reasoning process into steps
 - Decompose the problem into simpler sub-problems
 - Solve simpler sub-problems

Least-to-Most Prompting

Question

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?

↓
Question
Decomposition

Q: How many cars are in the parking lot after the first car arrives?
Q: How many cars are in the parking lot after the second car arrives?

↓
Subproblem
Solving

Q: How many cars are in the parking lot after the first car arrives?

A: There are originally 3 cars. After the first car arrives, we have $3 + 1 = 4$ cars.

Q: How many cars are in the parking lot after the second car arrives?

A: After the first car arrives, we have $3 + 1 = 4$ cars.
After the second car arrives, we have $4 + 1 = 5$ cars.

Least-to-Most Prompting

Least-to-most prompting (solving stage)

Q: “think, machine”

A: The last letter of “think” is “k”. The last letter of “machine” is “e”. Concatenating “k”, “e” leads to “ke”. So, “think, machine” outputs “ke”.

Q: “think, machine, learning”

A: “think, machine” outputs “ke”. The last letter of “learning” is “g”. Concatenating “ke”, “g” leads to “keg”. So, “think, machine, learning” outputs “keg”.


Table 2: A test case of least-to-most prompting for the last-letter-concatenation task. Generated with code-davinci-002 in GPT-3. The prompt context is shown on the right column of Table 1.

Method	L = 4	L = 6	L = 8	L = 10	L = 12
Standard prompting	0.0	0.0	0.0	0.0	0.0
Chain-of-Thought	89.4	75.0	51.8	39.8	33.6
Least-to-Most	94.0	88.4	83.0	76.4	74.0


Complex Few-shot Reasoning with LLMs: Key Techniques




Reasoning Elaboration: Spell out the reasoning process before generating the answer



Tool Augmentation



Structured Generation



Feedback



Memory

Instructions

- Answer the following question: Does water exist in a liquid state at room temperature?
 - Answer the following question: Does water exist in a liquid state at room temperature? Choose yes or no.
- Write a summary of the history of computer science
 - Summarize the history of computer science, highlighting key developments.
 - Provide a summary of the history of computer science in this format:
 - ❖ Inventors and innovations
 - ❖ Evolution of programming languages
 - ❖ Impact of internet
 - Summarize the history of computer science in no more than 200 words, keeping the language simple and easy to understand
 - Assume you're writing for a tech blog aimed at beginners. Summarize the history of computer science in a way that sparks interest in the subject.

Same end goal – different phrasing

- Implicit vs. Explicit
- Abstract vs. Concrete

Reframing Instructional Prompts to GPTk's Language (Mishra et al., 2022)

Raw task definitions and their reframed counterpart

Raw Task: Craft a question which requires commonsense to be answered. Based on the given context, craft a common-sense question, especially those that are LONG, INTERESTING, and COMPLEX. The goal is to write questions that are easy for humans and hard for AI machines! To create such questions, here are some suggestions: A. What may (or may not) be the plausible reason for an event? B. What may (or may not) happen before (or after, or during) an event? C. What may (or may not) be a plausible fact about someone (or something)? D. What may (or may not) happen if an event happens (or did not happen)? You can also create other types of questions.

Input: Context:<> **Expected Output:** Question:<>

Reframed Task: Use 'what may happen', 'will ...?', 'why might', 'what may have caused', 'what may be true about', 'what is probably true about', 'what must' and similar phrases in your question based on the input context.

Input: Context:<> **Expected Output:** Question:<>

Instructions

Raw Task:... *What is the type of the answer corresponding to the given question? Number, Date, or Span?...*

Input: Passage: <>. Question: <> **Expected Output:** <Number/Date/Span> ...

Reframed Task:... *What is the type of the answer corresponding to the given question? Number, Date, or Span?...*

Input: Passage: <> Question: <> *Answer either Number, Date or Span?* **Expected Out-**
put:<Number/Date/Span>

Instructions




Figure 2: Across a variety of model sizes, **reframed prompts** consistently show considerable performance gain over **raw task instructions (no reframing)** in a few-shot learning setup. Since fine-tuning GPT3 is


[Mishra, Swaroop, Daniel Khashabi, Chitta Baral, Yejin Choi, and Hannaneh Hajishirzi. "Reframing Instructional Prompts to GPTk's Language." *ACL 2022 Findings*.](#)

- Explicit is better than implicit
- Be aware of instances where you might be expecting the model to read your mind!


Complex Few-shot Reasoning with LLMs: Key Techniques




Reasoning Elaboration: Spell out the reasoning process before generating the answer




Instructions Provide explicit instructions to the LLM, capture all expectations



Tool Augmentation



Structured Generation



Feedback



Memory

Augmenting LLMs with Tools

- Typical language generation setup: LLMs handle *everything*
 - There are originally 3 cars. 2 more cars arrive. $3 + 2 = 5$.
 - What has to be done?
 - Add $3 + 2$
 - How to do it?
 - Perform the calculation and produce the answer = 5
- Leverage external tools to separate 'what' needs to be done from 'how' to do it.
- Calculator as a tool (T):
 - There are originally 3 cars.
 - 2 more cars arrive.
 - Number of cars after the first car comes = calculator(3, 1) = 4.
 - Number of cars after the second car comes = calculator(4, 1) = 5.
 - The answer is 5.

React (Yao et al. 2022)

(1) Hotspot QA

Question: Aside from the Apple Remote, what other device can control the program Apple Remote was originally designed to interact with?

(1a) Standard

Answer: iPod

(1c) Act-Only

Act 1: `Search[Apple Remote]`
Obs 1: The Apple Remote is a remote control ...

Act 2: `Search[Front Row]`
Obs 2: Could not find [Front Row]. Similar: ...

Act 3: `Search[Front Row (software)]`

Obs 3: Front Row is a discontinued media center software ...

Act 4: `Finish[yes]`

(1d) ReAct (Reason + Act)

Thought 1: I need to search Apple Remote and find the program it was originally designed to interact with.

Act 1: `Search[Apple Remote]`

Obs 1: The Apple Remote is a remote control introduced in October 2005 by Apple ... originally designed to control the **Front Row media center** program ...

Thought 2: Apple Remote was originally designed to control the Front Row media center program. I need to search **Front Row** next and find what other device can control it.

Act 2: `Search[Front Row]`

Obs 2: Could not find [Front Row]. Similar: ['Front Row Seat to Earth', 'Front Row Motorsports', 'Front Row (software)', ...]

Thought 3: Front Row is not found. I need to search **Front Row (software)**.

Act 3: `Search[Front Row (software)]`

Obs 3: Front Row is a discontinued media center software ...

Thought 4: Front Row (software) is controlled by an **Apple Remote** or the **keyboard function keys**. So the answer is keyboard function keys.

Act 4: `Finish[keyboard function keys]`



React (Yao et al. 2022)





Figure 3: Scaling results for prompting and finetuning on HotPotQA with ReAct (ours) and baselines.

Augmenting LLMs with Tools

- Obvious use cases
 - **Real-time information:** Stock market updates, Temperature monitoring, Traffic reports.
 - **Specialization:** Solving Ordinary Differential Equations (ODEs)
 - **Multimodal:** Performing clicks, Generating image captions.
- Future Directions
 - Combine multiple tools within the same language model framework for increased functionality and adaptability.
 - [\[2304.09842\] Chameleon: Plug-and-Play Compositional Reasoning with Large Language Models](#)
 - [\[2304.08354\] Tool Learning with Foundation Models](#)
 - [\[2302.04761\] Toolformer: Language Models Can Teach Themselves to Use Tools](#)
 - [\[2303.09014\] ART: Automatic multi-step reasoning and tool-use for large language models](#)

Complex Few-shot Reasoning with LLMs: Key Techniques




Reasoning Elaboration: Spell out the reasoning process before generating the answer




Instructions Provide explicit instructions to the LLM, capture all expectations



Tool Augmentation: Enhance LLMs for specialized tasks



Structured Generation




Feedback



Memory

Structured Commonsense Reasoning

- Natural language input (e.g., scenario)
- Structured output (e.g., plan graph, reasoning graph)




Leveraging Language Models for Structured commonsense Reasoning

- Need to generate a graph but ... language models can only generate strings
- Workaround
 - *Flatten* the graph as a string
 - Train a seq2seq model

Neural Language Modeling for Contextualized Temporal Graph Generation
Aman Madaan, Yiming Yang


proScript: Partially Ordered Scripts Generation
Keisuke Sakuruchi,¹ Chandra Bhagavatula,¹ Ronan Le Bras,¹
Niket Tandon,¹ Peter Clark,² Yejin Choi^{1,2}
¹Allen Institute for Artificial Intelligence
²Paul G. Allen School of Computer Science & Engineering, University of Washington

Goal: Bake a cake



```
.....  
:"find recipe" -> "gather ingredients";  
:"gather ingredients" -> "mix ingredients";  
:"gather ingredients" -> "preheat oven at 375F";  
:"find recipe" -> "preheat oven at 375F";  
:"preheat oven at 375F" -> "put the cake batter into oven";  
:"mix ingredients" -> "put the cake batter into oven";  
:"put the cake batter into oven" -> "take the cake out after 30 min"  
.....
```

Intermediate Representation




Recovered Graph

Leveraging Language Models for Structured commonsense Reasoning

- Issues with the workaround

```
"find recipe" → "gather ingredients";  
"gather ingredients" → "mix ingredients";  
"find recipe" → "preheat oven at 375F";  
"preheat oven at 375F" → "put the cake batter into oven";  
"mix ingredients" → "put the cake batter into oven";  
"put the cake batter into oven" → "take the cake out after 30 min"
```



- We want **structures, not strings**


!? Are the two
mix ingredients the same?

Language Models of Code Are Few-shot Commonsense Learners (Madaan et al. 2022)

```
class BakeACake:  
    def __init__(self) -> None:  
        self.find_recipe = Node()  
        self.gather_ingredients = Node()  
        self.mix_ingredients = Node()  
        self.find_recipe = Node()  
        self.preheat_oven_at_375f = Node()  
        self.put_cake_batter_into_oven = Node()  
        self.take_cake_out_after_30_min = Node()
```

Goal: Bake
a cake

```
    self.find_recipe.children = [self.gather_ingredients,  
    self.preheat_oven_at_375f]  
    self.gather_ingredients.children = [self.mix_ingredients]  
    self.mix_ingredients.children = [self.put_cake_batter_into_oven]  
    self.preheat_oven_at_375f.children =  
    [self.put_cake_batter_into_oven]  
        self.put_cake_batter_into_oven.children =  
    [self.take_cake_out_after_30_min]
```



Intermediate Representation

Recovered Graph


Language Models of Code Are Few-shot Commonsense Learners (Madaan et al. 2022)

Goal: Bake a cake



```
■"find recipe" -> "gather ingredients";
  "gather ingredients" -> "mix ingredients";
■"gather ingredients" -> "preheat oven at 375F";
  "find recipe" -> "preheat oven at 375F";
■"preheat oven at 375F" -> "put the cake batter into oven";
  "mix ingredients" -> "put the cake batter into oven";
■"put the cake batter into oven" -> "take the cake out after 30 min"
```

Intermediate Representation



Recovered Graph


Goal: Bake a cake



```
class BakeACake:
    def __init__(self) -> None:
        self.find_recipe = Node()
        self.gather_ingredients = Node()
        self.mix_ingredients = Node()
        self.find_recipe = Node()
        self.preheat_oven_at_375f = Node()
        self.put_cake_batter_into_oven = Node()
        self.take_cake_out_after_30_min = Node()


        self.find_recipe.children = [self.gather_ingredients,
        self.preheat_oven_at_375f]
        self.gather_ingredients.children = [self.mix_ingredients]
        self.mix_ingredients.children = [self.put_cake_batter_into_oven]
        self.preheat_oven_at_375f.children =
        [self.put_cake_batter_into_oven]
        self.put_cake_batter_into_oven.children =
        [self.take_cake_out_after_30_min]
```

Intermediate Representation



Recovered Graph

Script Generation Results on ProScript



- Olivia has \$23. She bought five bagels for \$3 each. How much money does she have left?

Olivia had 23 dollars. 5 bagels for 3 dollars each will be dollars. So she has dollars left.

```
def solution():
    money_initial = 23
    bagels = 5
    bagel_cost = 3
    money_spent = bagels * bagel_cost
    money_left = money_initial - money_spent
    result = money_left
    return result
```




Comparison with CoT:

- The language model is responsible for generating a high-level plan that is executed to derive the answer
- The results are obtained after running the program

Improves Solve Rate for Multiple Maths Reasoning Tasks

Solve Rate



Why should code help?

```
class BakeACake:  
    def __init__(self) -> None:  
        self.find_recipe = Node()  
        self.gather_ingredients = Node()  
        self.mix_ingredients = Node()  
        self.find_recipe = Node()  
        self.preheat_oven_at_375f = Node()  
        self.put_cake_batter_into_oven = Node()  
        self.take_cake_out_after_30_min = Node()  
  
        self.find_recipe.children = [self.gather_ingredients,  
self.preheat_oven_at_375f]  
        self.gather_ingredients.children = [self.mix_ingredients]  
        self.mix_ingredients.children = [self.put_cake_batter_into_oven]  
        self.preheat_oven_at_375f.children =  
[self.put_cake_batter_into_oven]  
        self.put_cake_batter_into_oven.children =  
[self.take_cake_out_after_30_min]
```

Structured Generation



Scott Condron  
@_ScottCondron

...


Building a classifier in 2023

Use @OpenAI's new function calling API to define the possible outputs and then use the "input" argument it returns as the classification

```
def classify(input_string: str) -> str:
    functions = [
        {
            "name": "print_sentiment",
            "description": "A function that prints the given sentiment",
            "parameters": {
                "type": "object",
                "properties": {
                    "sentiment": {
                        "type": "string",
                        "enum": ["positive", "negative", "neutral"],
                        "description": "The sentiment."
                    },
                    "required": ["sentiment"]
                }
            }
        }
    ]
    messages = [{"role": "user", "content": input_string}]
    response = openai.ChatCompletion.create(
        model="gpt-3.5-turbo-0613",
        messages=messages,
        functions=functions,
        function_call={"name": "print_sentiment"},
    )
    function_call = response.choices[0].message["function_call"]
    argument = json.loads(function_call["arguments"])
    return argument
```

https://twitter.com/_ScottCondron/status/1670827747684364288

Complex Few-shot Reasoning with LLMs: Key Techniques





Reasoning Elaboration: Spell out the reasoning process before generating the answer



Instructions Provide explicit instructions to the LLM, capture all expectations



Tool Augmentation: Enhance LLMs for specialized tasks



Structured Generation: Use structure to guide model



Memory

Feedback

- Humans typically generate content iteratively:
 - Generate → Feedback → Refine
- LLMs, with the right prompt, have recently shown promise at mimicking this process
 - Generate an output
 - Gather feedback
 - Improve output with feedback, repeat

Self-Refine (Madaan et al. 2023)

USER Give me an acronym for a paper titled "Politeness transfer: A tag and generate approach"

Make it easy to pronounce.
Relevant to the title.

ASSISTANT PTAGA (Puh-TAG-uh)

USER That is not a great acronym either by relevance or pronunciation

ASSISTANT My apologies for that. Here's another acronym: P-TAG (Pee-Tag)

USER Same problem!

ASSISTANT Apologies for the confusion. Here's another attempt: TRAPG (Transferring Respect and Politeness: Generate Approach) which can be pronounced as "trap-gee".

Mode Beta

Model

Temperature 0.7

Maximum length 256

Top P 1

Frequency penalty 0


Presence penalty 0

GPT-3.5/4 is good!
But often requires a
few rounds of talking
through the problem




Let LLM Self-Refine



Generate Acronym for: *Politeness transfer: A tag and generate approach*



PTATAGA



POLTAG



Relevance of PTATAGA: The acronym includes the first letters of each major word in the title but doesn't fully capture the meaning of the title.
<3/5>

Memorability of PTATAGA: The acronym is not the most memorable.
<2/5>

Pronunciation of PTATAGA: The acronym is pronounceable but not very fluid. <3/5>

Relevance of POLTAG: A tag and generate approach:
The acronym POLTAG is a good representation of the title as it combines the first few letters of Politeness and Tag.
<4/5>

Memorability of POLTAG: The acronym POLTAG is easy to remember as it is short, simple, and has a clear structure. <4/5>

Pronunciation of POLTAG: The acronym is easily pronounced as it follows the common structure of acronyms. <5/5>

STOP

Results


Task	GPT-3.5		ChatGPT		GPT-4	
	Base	+SELF-REFINE	Base	+SELF-REFINE	Base	+SELF-REFINE
Sentiment Reversal	8.8	30.4 (\uparrow 21.6)	11.4	43.2 (\uparrow 31.8)	3.8	36.2 (\uparrow 32.4)
Dialogue Response	36.4	63.6 (\uparrow 27.2)	40.1	59.9 (\uparrow 19.8)	25.4	74.6 (\uparrow 49.2)
Code Optimization	14.8	23.0 (\uparrow 8.2)	23.9	27.5 (\uparrow 3.6)	27.3	36.0 (\uparrow 8.7)
Code Readability	37.4	51.3 (\uparrow 13.9)	27.7	63.1 (\uparrow 35.4)	27.4	56.2 (\uparrow 28.8)
Math Reasoning	64.1	64.1 (0)	74.8	75.0 (\uparrow 0.2)	92.9	93.1 (\uparrow 0.2)
Acronym Generation	41.6	56.4 (\uparrow 14.8)	27.2	37.2 (\uparrow 10.0)	30.4	56.0 (\uparrow 25.6)
Constrained Generation	28.0	37.0 (\uparrow 9.0)	44.0	67.0 (\uparrow 23.0)	15.0	45.0 (\uparrow 30.0)

Table 1: SELF-REFINE results on various tasks using GPT-3.5, ChatGPT, and GPT-4 as base LLM. SELF-REFINE consistently improves LLM. Metrics used for these tasks are defined in Section 3.2.


- Larger gains where feedback is obvious + easy
- Lots of room for improvement!

Future Directions: Feedback + Planning

- Monte-Carlo Tree Search + Self-Refine
 - Extend Self-Refine Beyond Linear Search
- Other exciting work:
 - Tree-of-thought
 - Reflexion
 -



Complex Few-shot Reasoning with LLMs: Key Techniques




Reasoning Elaboration: Spell out the reasoning process before generating the answer




Instructions Provide explicit instructions to the LLM, capture all expectations




Tool Augmentation: Enhance LLMs for specialized tasks



Structured Generation: Use structure to guide model



Feedback: Refine model outputs during inference




Memory

- Standard few-shot prompting setup
 - Prompt $P: [X_i, Y_i]$
 - Test examples X_1, X_2
 - $P + X_1 \rightarrow Y'_1$
 - $X_2 \rightarrow Y'_2$
- Idea: maintain a memory of examples seen so far, and any feedback


MemPrompt (Madaan et al. 2022)

Models repeat mistakes



GPT-4



MemPrompt




Workflow

- Step 1: User asks a **question**
- Step 2: Check if the same **question** has been asked before, and a **clarification** is present in memory
 - Step 2.1: If a **clarification** is present, add **question + clarification**
 - Step 2.2: If not, just ask a **question**
- Step 3: Model generates an **answer**
- Step 4: Take **clarification** on **answer** if needed, add **clarification** to memory

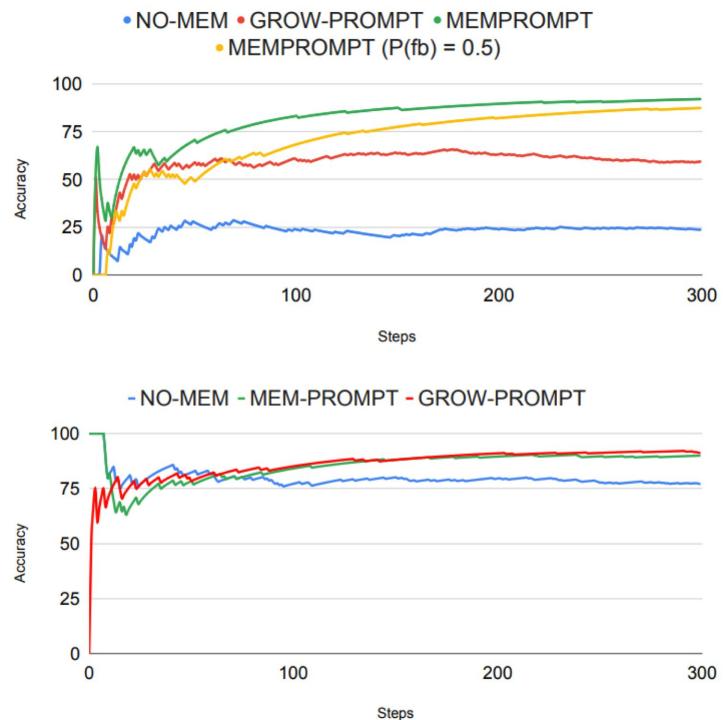
Question	Feedback
A word pronounced as fellow ?	I want a word that sounds similar!
What is dissimilar to delicious ?	Give me the reverse of delicious
What is a word like great ?	Wrong! I want something similar ✓
How do I use melancholy ?	No...I wanted a sample sentence
What is on the lines of pretty ?	I was looking for a similar word
Could you expand on browser ?	I actually wanted a definition



Memprompt Results



Ethical Reasoning with Categorical Feedback




Lexical tasks (top) and word scrambling tasks (bottom)

Memory for Few-shot Prompting

- Simpler but effective variants
 - KATE: Store the training set in a database, retrieve most relevant examples on the fly

Complex Few-shot Reasoning with LLMs: Key Techniques




Reasoning Elaboration: Spell out the reasoning process before generating the answer




Instructions Provide explicit instructions to the LLM, capture all expectations



Tool Augmentation: Enhance LLMs for specialized tasks



Structured Generation: Use structure to guide model




Feedback: Refine model outputs during inference



Memory: maintain a history of interactions with LLM


Future Directions

- LLMs are getting better at following instructions, exciting new possibilities
 - Planning, Search – resurgence of classical AI techniques
 - Generating actionable feedback from tools
 - Storing a memory of interactions across multiple users
 - Enhanced Structured Generation Capabilities



Complex Few-shot Reasoning with LLMs: Key Techniques

Questions?





Reasoning Elaboration: Spell out the reasoning process before generating the answer



Instructions Provide explicit instructions to the LLM, capture all expectations



Tool Augmentation: Enhance LLMs for specialized tasks



Structured Generation: Use structure to guide model




Feedback: Refine model outputs during inference



Memory: maintain a history of interactions with LLM

Appendix

Text and Patterns: For Effective Chain of Thought, It Takes Two to Tango



Aman Madaan* and **Amir Yazdanbakhsh***
Carnegie Mellon University *Google Research, Brain Team
amadaan@cs.cmu.edu, ayazdan@google.com
(Equal Contribution)


[Paper](#)

What makes the chain of thought prompting so effective?


- *The thought makes the model think about the problem?*
- *The thought helps the model learn better*
- *The thought serves as an additional example of the task*
- *The thought helps the model remind of the task*
- *The thought helps extract relevant information for solving the task*

What makes chain of thought prompting so effective

Text: Extracting common sense, copy mechanism



Patterns: Task understanding, final answer generation



Approach

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?

T: There are originally 3 cars. 2 more cars arrive. $3 + 2 = 5$.

A: The answer is 5 cars.

Symbols

Patterns

Text

- ***Counterfactual prompting:***

- Change one *knob* at a time (symbol, patterns, text)



What if? prompting (counterfactual prompting)

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?

Thought (T): There are originally 3 cars. 2 more cars arrive. $3 + 2 = 5$.

A: The answer is 5 cars.

What if we don't have actual numbers?

What if the prompt is misleading?

Q: If there are α cars in the parking lot and β more cars arrive, how many cars are in the parking lot?

Thought (T): There are originally α cars. β more cars arrive. $\alpha + \beta = \lambda$.


A: The answer is λ cars.

Q: If there are 3 cars in the parking lot and 2 more cars arrive, how many cars are in the parking lot?


Thought (T): There are originally 3 cars. 2 more cars arrive. $3 + 2 = 7$.

A: The answer is 5 cars.

Evaluating counterfactual prompts



Empirical results: difference in final outcome



Attention patterns: difference in mechanism

Complex Reasoning with LLMs

- Complex Reasoning with LLMs
- Few-shot prompting
 - Why does it work?
- Augmenting Complex Reasoning
 - **Better Chain-of-thought**
 - Least2Most, Algorithmic Prompt, Self-ask
 - Programs
 - PaL, ReACT
 - Memory
 - Memprompt, KATE
 - Structures
 - Tree-of-thought
 - Sampling
 - Re-prompting

Complex Reasoning with LLMs

- Few-shot prompting
 - Why does it work?
- Class of Techniques for Few-shot Prompting
 - Decomposition
 - Least2Most, Algorithmic Prompt, Self-ask, Complex Prompting
 - Code-Assisted Reasoning
 - PaL, PoT, ReACT
 - Structures
 - - Memory: Memprompt, KATE
 - Trees: Tree-of-thought
 - Sampling
 - Re-prompting, Self-Refine